



- ▶ Hiwi-Job-Angebot DFKI (siehe read.mi)
- ▶ Hinweis: QIS-Anleitung (siehe read.mi)
- ▶ Slides letztes Kapitel (Feedback) nachholen



# Applications of Artificial Intelligence

– Winter Term 21/22 –

## Chapter 04

# Natural Language Processing


Prof. Dr. Adrian Ulges

RheinMain University of Applied Sciences



1. Fundamentals
2. Preprocessing and Heuristics
3. Statistical Word Similarity
4. Syntactical Word Similarity

## Recommended Read [3]




# Natural Language Processing: Some Tasks<sup>1</sup>



## Part-of-Speech Tagging

Abraham Lincoln was the 16th President of the United States. He was **shot** by stage actor John Wilkes Booth. It was in his **interest** ...



## Coreference Resolution

**Abraham Lincoln** was the 16th President of the United States. **He** was assassinated by stage actor John Wilkes Booth on April 14, 1865.

## Named Entity Recognition

**Abraham Lincoln** was the 16th President of the **United States**. He was assassinated by stage actor **John Wilkes Booth** on **April 14, 1865**.



## Information Extraction


**Abraham Lincoln** was the 16th President of the United States. He was **assassinated** by stage actor **John Wilkes Booth** on April 14, 1865.



## Question Answering

> "Who killed Abe Lincoln?"

Abraham Lincoln was the 16th President of the United States. He was assassinated by stage actor John Wilkes Booth on April 14, 1865.



## Goal-oriented Dialog Understanding

**user**

> What's a good diner nearby?


> How far is this?

> "How about Italian ones?"

**bot**

> How about the Taj Mahal?

> 500 m



<sup>1</sup>The Natural Language Decathlon: <https://github.com/salesforce/decaNLP>

# Natural Language Processing: Basics



## Morphology

- ▶ analyzes the structure and parts of words (walk - ed).

## Syntax (*greek "order together"*)

- ▶ How are words ordered? What is a sentence's structure?

## Semantics



- ▶ What is the **meaning** of a word/sentence?
- ▶ **Goal** of many NLP techniques
  - ▶ **classification** of sentences / questions
  - ▶ **extraction** of entities / facts

# Text → Semantics?



Is the Eiffeltower higher than Notre Dame ?

The Eiffeltower is an iron lattice tower on the Champ de Mars.  
With its **height of 324 meters**, it is the **tallest structure** in Paris.



Transferring natural language into a formal description is **not entirely solved** so far.

## Example: The Research Project **Read The Web**<sup>2</sup>

cracking\_new\_year is a monarch

1045 30-mar-2017

96.3 🐼🐼

tuna is a type of large predatory\_fish

1046 02-apr-2017

start\_search\_field is a research\_project

1046 02-apr-2017

98.8 🐼🐼

tampa\_bay\_devil\_rays is headquartered\_in the state or province new\_york

1050 19-apr-2017

adrianne\_curry is a fashion\_model

1045 30-mar-2017

99.8 🐼🐼

man is an animal that can develop\_disability

1046 02-apr-2017

<sup>2</sup><http://rtw.ml.cmu.edu/rtw/>



1. Fundamentals
2. Preprocessing and Heuristics
3. Statistical Word Similarity
4. Syntactical Word Similarity



# Text Preprocessing



Often, the first step in NLP systems is preprocessing text.  
Here, two possible steps are:

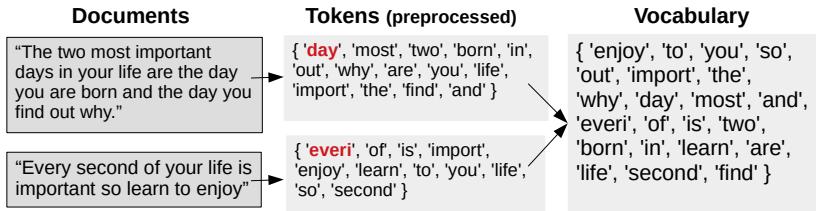
## 1. Segmentation

- ▶ Segmenting text into **sentences**
- ▶ Segmenting sentences into **words/terms/tokens**  
(*"house"*, *"rock 'n roll"*, *"#nlproc"*,  
*"getStringFromObject()"*?)
- ▶ This is often done using **regular expressions**.

## 2. Normalization

- ▶ reducing words to their base form (*aka 'lemma'*).
- ▶ e.g., lemmatization, stemming, compound splitting.

# Text Preprocessing (cont'd)



## 3. Vocabulary

- ▶ Collect all tokens from the corpus in a **vocabulary**  $\{t_1, \dots, t_m\}$ .
- ▶ Usually, we filter tokens that are too frequent, too infrequent, numbers, special tokens, etc.

# Regular Expressions



- ▶ ... match patterns in strings, using a compact formal language.
- ▶ Symbols in []-brackets: logical OR.

|               |                          |                                    |
|---------------|--------------------------|------------------------------------|
| [Ww]ood       | Wood / wood              | <b>Wood</b> stock is small indeed. |
| [0 - 9] or \d | some digit               | Beverly Hills <b>9</b> 0210.       |
| yours mine    | alternative: pipe symbol | Brace <b>y</b> ourself.            |
| [a-z]         | Kleinbuchstaben          | <b>W</b> oodstock is small indeed. |

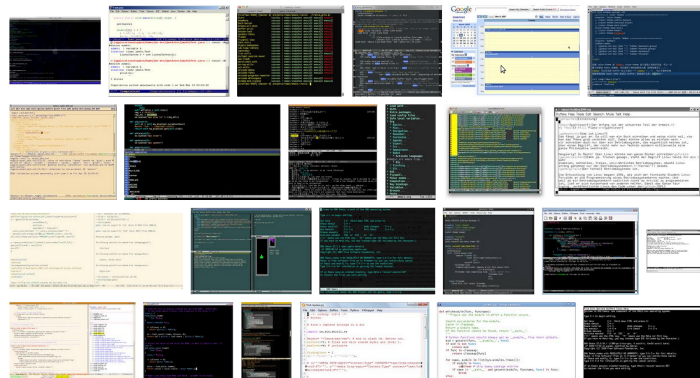
- ▶ Negation: [^X] means “not X”.

|        |  |                       |
|--------|--|-----------------------|
| [^A-Z] | anything but capital letters             | <b>W</b> oodstock.    |
| [^sS]  | anything but s or S                      | <b>S</b> ure you can. |
| a^b    | a^b (no negation, since not in brackets) | Sure you can.         |

- ▶ Other Options:

|         |                         |   |
|---------|-------------------------|---|
| colou?r | u is optional           | <b>colors</b> and <b>colours</b> .      |
| ye*ah   | 0 or more occurrences   | <b>yah!</b> <b>yeah!</b> <b>yeeeah!</b> |
| ye+ah   | 1 or more occurrences   | yah! <b>yeah!</b> <b>yeeeah!</b>        |
| beg.n   | any symbol except \n    | begn <b>begin began beg3n</b>           |
| .*      | any sequence of symbols | <b>abcdefghijklmnop</b>                 |
| yeah\$  | yeah at end of string   | yeah yeah <b>yeah</b>                   |

# Regular Expressions in Python



## Applications (e.g. in QA)

- ▶ **Segmenting**, rule-based **recognition** of dates, prices, ...
- ▶ **Features** for QA's machine learning components  
(*found abbreviation?* → *maybe question of type "definition"*).

# Rule-based Inference with RegExp: Hearst Patterns [1] \*

|                                     |   |
|-------------------------------------|---|
| NP such as {NP,}* {(and/or)} NP     | → Football teams, such as the Eagles or the Patriots, ... |
| NP, especially {NP,}* {(and/or)} NP | → I really like animals, especially birds of prey.        |

- ▶ Lexico-syntactic patterns to recognize hyponymy.
- ▶ <NP>s are so-called **noun phrases**, such as *birds of prey* or *the Eagles* (more later).
- ▶ Typical for pattern-based approaches:  
high precision 😊, low recall ☹.
- ▶ Some Patterns in Python<sup>3</sup> (47 patterns total) are below.

```
1 [
2   '(<NP>\w+ (, )?such as (<NP>\w+ ?(, )?(and / or )?)+',
3   '(<NP>\w+ (, )?as (<NP>\w+ ?(, )?(and / or )?)+',
4   '((<NP>\w+ ?(, )?)+(and / or )? other <NP>\w+)',
5 ]
```

<sup>3</sup>[https://github.com/mmichelsonIF/hearst\\_patterns\\_python](https://github.com/mmichelsonIF/hearst_patterns_python)

# Words: Reduction to Base Forms



- ▶ Words are the smallest **meaningful units** of a language.
- ▶ Words come in a **baseform** and multiple **variants**.
- ▶ Often, we **normalize/canonicalize** words to their base form.

## Canonicalization

- ▶ **from flexion form** to base form

*Hauses* → *Haus*

*went* → *go*

- ▶ **derivational forms** to base form

*Bücherei* → *Buch*

- ▶ from **compounds** into components (*especially for German*)

*Haustür* → *Haus* + *Tür*

*Osterei* → *Ostern* + *Ei*

*Malerei* ↗ *Maler* + *Ei*



## Lemmatization

- ▶ flexion form → base form (*above*, e.g. housing → house).
- ▶ A simple form of lemmatization is **stemming**: simply truncate a word suffix (housing → hous).

## Stemming: Approaches

### 1. Lexical Methods

- ▶ store base form in dictionaries
- ▶ creation and maintenance of dictionary costly ☹
- ▶ interesting for languages with **strong flexion** (e.g., *German*)

### 2. Rule-based Methods

- ▶ use rules to change / remove suffixes
- ▶ example: **Porter Stemmer** for English [3]

```
ing → - : walking → walk  
ies → i : ponies → poni
```

# Parts of Speech (*dt. Wortarten*)



The four most important “**open**” parts of speech.

## Nouns (*dt. Nomen*)

- ▶ ... refer to concrete things (ship), abstract things (bandwidth), actions (belief), events (inspection), ...
- ▶ ... can be accompanied by determiners (some bandwidth) und possessives (his ship).
- ▶ special case: **proper nouns** (*dt. Eigennamen*), e.g. Colorado.

## Verbs

- ▶ ... refer to **activities** and **actions**
- ▶ ... are changed through **flexion**: eat, eats, ate, eaten

## Adjectives

- ▶ ... refer to **properties** of **nouns** (red ship)

## Adverbs

- ▶ ... refer to **properties** of **verbs** (he walked slowly)



# How many Parts of Speech are there in English?



45 (in the famous “Penn Treebank”)

| Tag   | Description          | Example               | Tag  | Description          | Example              |
|-------|----------------------|-----------------------|------|----------------------|----------------------|
| CC    | coordin. conjunction | <i>and, but, or</i>   | SYM  | symbol               | <i>+, %, &amp;</i>   |
| CD    | cardinal number      | <i>one, two</i>       | TO   | “to”                 | <i>to</i>            |
| DT    | determiner           | <i>a, the</i>         | UH   | interjection         | <i>ah, oops</i>      |
| EX    | existential ‘there’  | <i>there</i>          | VB   | verb base form       | <i>eat</i>           |
| FW    | foreign word         | <i>mea culpa</i>      | VBD  | verb past tense      | <i>ate</i>           |
| IN    | preposition/sub-conj | <i>of, in, by</i>     | VBG  | verb gerund          | <i>eating</i>        |
| JJ    | adjective            | <i>yellow</i>         | VBN  | verb past participle | <i>eaten</i>         |
| JJR   | adj., comparative    | <i>bigger</i>         | VBP  | verb non-3sg pres    | <i>eat</i>           |
| JJS   | adj., superlative    | <i>wildest</i>        | VBZ  | verb 3sg pres        | <i>eats</i>          |
| LS    | list item marker     | <i>1, 2, One</i>      | WDT  | wh-determiner        | <i>which, that</i>   |
| MD    | modal                | <i>can, should</i>    | WP   | wh-pronoun           | <i>what, who</i>     |
| NN    | noun, sing. or mass  | <i>llama</i>          | WP\$ | possessive wh-       | <i>whose</i>         |
| NNS   | noun, plural         | <i>llamas</i>         | WRB  | wh-adverb            | <i>how, where</i>    |
| NNP   | proper noun, sing.   | <i>IBM</i>            | \$   | dollar sign          | <i>\$</i>            |
| NNPS  | proper noun, plural  | <i>Carolinas</i>      | #    | pound sign           | <i>#</i>             |
| PDT   | predeterminer        | <i>all, both</i>      | “    | left quote           | <i>‘ or “</i>        |
| POS   | possessive ending    | <i>’s</i>             | ”    | right quote          | <i>’ or ”</i>        |
| PRP   | personal pronoun     | <i>I, you, he</i>     | (    | left parenthesis     | <i>[, (, {, &lt;</i> |
| PRP\$ | possessive pronoun   | <i>your, one’s</i>    | )    | right parenthesis    | <i>], ), }, &gt;</i> |
| RB    | adverb               | <i>quickly, never</i> | ,    | comma                | <i>,</i>             |
| RBR   | adverb, comparative  | <i>faster</i>         | .    | sentence-final punc  | <i>. ! ?</i>         |
| RBS   | adverb, superlative  | <i>fastest</i>        | :    | mid-sentence punc    | <i>; ; ... --</i>    |
| RP    | particle             | <i>up, off</i>        |      |                      |                      |

# Part-of-Speech (POS) Tagging



Parts of Speech (at least in English) are (highly) ambiguous!

book that flight. vs. hand me this book.  
book that flight. vs. I know that you stink.

- ▶ Most words (86%) have only one part of speech, but some **frequent** words have **multiple** POSs (that, back, set, ...).
- ▶ English texts contain  $\approx$  **60% ambiguous tokens**.

## *Part-of-speech-Tagging (POS-Tagging)*

= given a text, determine each token's POS.

- ▶ Core step for **many applications**: recognizing entities, parsing, key word detection ...
- ▶ 92% accuracy by simple baseline (choosing most frequent POS per word), 97% accuracy by **machine learning** (*later*).

# POS-Tagging in Python



```
1 import nltk
2
3 text = "We watched the sunset over the castle on the hill."
4
5 tokens = nltk.word_tokenize(text)
6
7 result = nltk.pos_tag(tokens)
8
9 print(result)
10
11 > [('We', 'PRP'),           # OK (personal pronoun)
12    ('watched', 'VBD'),      # OK (verb past tense)
13    ('the', 'DT'),          # OK (determiner)
14    ('sunset', 'NN'),       # OK (noun singular)
15    ('over', 'IN'),         # OK (preposition)
16    ('the', 'DT'),          # OK (determiner)
17    ('castle', 'NN'),       # OK (noun singular)
18    ('on', 'IN'),           # OK (preposition)
19    ('the', 'DT'),          # OK (determiner)
20    ('hill', 'NN'),         # OK (noun singular)
21    ('.', '.')]

```



1. Fundamentals
2. Preprocessing and Heuristics
3. Statistical Word Similarity
4. Syntactical Word Similarity



Q: Which **insurance** company is based in **Wiesbaden**?

A: The R+V, a big **insurer**, resides in **Wiesbaden**.

- ▶ For QA (and many other applications), it is interesting to find **similar words**.
- ▶ Example: **query expansion**: (laptop → notebook).
- ▶ There are several forms of similarity (*semantic, statistical, syntactic*).
- ▶ Sometimes, we can identify similar words by **syntactic** similarity (i.e., similarity of the word form).
- ▶ In other cases, this is not enough:

insurance → insurer ✓

laptop → notebook ✗



## Statistical Word Similarities

- ▶ ... **learn** word similarity from a **text corpus**.
- ▶ Words are similar if their **contexts** are similar!

*"You shall know a word by the company it keeps."*

(J.R.Firth (1957))

Example: What other words is "tesgüino" similar to?

*A bottle of **tesgüino** is on the table.*

***Tesgüino** makes you drunk.*

*We make **tesgüino** out of corn.*

# Statistical Word Similarity



There are two kinds of statistical similarity

- ▶ **first-order** (“collocations”):  $w_2$  is similar to  $w_1$ , if  $w_2$  appears **in  $w_1$ 's context** frequently (e.g. wrote  $\rightarrow$  book).
- ▶ **second-order**:  $w_2$  is similar to  $w_1$  if  $w_1$  and  $w_2$  tend to appear in **similar contexts** (e.g. vodka  $\rightarrow$  whisky).
- ▶ Context usually refers to a local window of  $\pm(1 - 8)$  words.

# Computing Statistical Word Similarities



- ▶ **Given:** vocabulary of  $n$  words + text corpus.
- ▶ **Goal:** compute an  $n \times n$ -matrix  $A$  (**term-term matrix**), such that  $A_{ij}$  is the similarity between  $w_i$  and  $w_j$ .
- ▶ Naive solution:

$$A_{ij} := \# \text{ occurrences of } w_j \text{ in the contexts of } w_i$$

- ▶ Why is this solution **bad**?

|         | wrote | book | ... | carbon | dioxide | ... | is  | of   |
|---------|-------|------|-----|--------|---------|-----|-----|------|
| wrote   | -     | 128  |     | 3      | 1       |     | 223 | 351  |
| book    |       | -    |     | 4      | 3       |     | 278 | 624  |
| ...     |       |      | ... |        |         |     |     |      |
| carbon  |       |      |     | -      | 46      |     | 98  | 75   |
| dioxide |       |      |     |        | -       |     | 44  | 52   |
| ...     |       |      |     |        |         | ... |     |      |
| is      |       |      |     |        |         |     | -   | 9748 |
| of      |       |      |     |        |         |     |     | -    |





## Positive Pointwise Mutual Information (PPMI) ([3], Chapter 15)

Let  $w_i$  be a word,  $w_j$  be a “context word”, and  $A$  the term-term matrix (*see above*).

We calculate

- ▶ The probability of  $w_i$  appearing in a **random local word pair**:

$$P(w_i) := \sum_j A_{ij} / \sum_{i',j'} A_{i'j'}$$

- ▶ The probability of  $w_j$  appearing as a **context term**:

$$P(w_j) := \sum_i A_{ij} / \sum_{i',j'} A_{i'j'}$$

- ▶ The probability that **both words** appear in the same pair:

$$P(w_i, w_j) := A_{ij} / \sum_{i',j'} A_{i'j'}$$



- ▶ We know (*statistics*): If both words are **independent**, then:

$$P(w_i, w_j) \approx P(w_i) \cdot P(w_j)$$

- ▶ We are interested in word pairs appearing **unusually frequently**, i.e.:

$$P(w_i, w_j) \gg P(w_i) \cdot P(w_j)$$

- ▶ This leads to the definition of **Pointwise Mutual Information (PMI)**:

$$PMI(w_i, w_j) := \log_2 \left( \frac{P(w_i, w_j)}{P(w_i) \cdot P(w_j)} \right)$$

# PMI Word Similarity



- ▶ The **higher** PMI, the more **similar** two words.
- ▶ **Example:** The top-10 and bottom-10 pairs from Wikipedia.

| word 1 | word 2    | count word 1 | count word 2 | count of co-occurrences | PMI            |
|--------|-----------|--------------|--------------|-------------------------|----------------|
| puerto | rico      | 1938         | 1311         | 1159                    | 10.0349081703  |
| hong   | kong      | 2438         | 2694         | 2205                    | 9.72831972408  |
| los    | angeles   | 3501         | 2808         | 2791                    | 9.56067615065  |
| carbon | dioxide   | 4265         | 1353         | 1032                    | 9.09852946116  |
| prize  | laureate  | 5131         | 1676         | 1210                    | 8.85870710982  |
| san    | francisco | 5237         | 2477         | 1779                    | 8.83305176711  |
| nobel  | prize     | 4098         | 5131         | 2498                    | 8.68948811416  |
| ice    | hockey    | 5607         | 3002         | 1933                    | 8.6555759741   |
| star   | trek      | 8264         | 1594         | 1489                    | 8.63974676575  |
| car    | driver    | 5578         | 2749         | 1384                    | 8.41470768304  |
| it     | the       | 283891       | 3293296      | 3347                    | -1.72037278119 |
| are    | of        | 234458       | 1761436      | 1019                    | -2.09254205335 |
| this   | the       | 199882       | 3293296      | 1211                    | -2.38612756961 |
| is     | of        | 565679       | 1761436      | 1562                    | -2.54614706831 |
| and    | of        | 1375396      | 1761436      | 2949                    | -2.79911817902 |
| a      | and       | 984442       | 1375396      | 1457                    | -2.92239510038 |
| in     | and       | 1187652      | 1375396      | 1537                    | -3.05660070757 |
| to     | and       | 1025659      | 1375396      | 1286                    | -3.08825363041 |
| to     | in        | 1025659      | 1187652      | 1066                    | -3.12911348956 |
| of     | and       | 1761436      | 1375396      | 1190                    | -3.70663100173 |

# From PMI to PPMI



Since we are not interested in negative values, we clip them to zero, and obtain the **Positive Pointwise Mutual Information (PPMI)**:

$$PPMI(w_i, w_j) := \max\left(\log_2\left(\frac{P(w_i, w_j)}{P(w_i) \cdot P(w_j)}\right), 0\right)$$

|         | wrote | book | ..  | carbon | dioxide | ..  | is  | of  |
|---------|-------|------|-----|--------|---------|-----|-----|-----|
| wrote   | -     | 2.9  |     | 0.5    | 0       |     | 0   | 0.1 |
| book    |       | -    |     | 0.2    | 0.2     |     | 0   | 0.1 |
| ...     |       |      | ... |        |         |     |     |     |
| carbon  |       |      |     | -      | 7.2     |     | 0.8 | 0   |
| dioxide |       |      |     |        | -       |     | 0.1 | 0   |
| ...     |       |      |     |        |         | ... |     |     |
| is      |       |      |     |        |         |     | -   | 0   |
| of      |       |      |     |        |         |     |     | -   |



## Remarks

- ▶ **Challenge:** very frequent words that co-occur spuriously sometimes get **extreme PPMI values**. Therefore, we **smooth** the term-term matrix  $A$  before computing PPMI by adding a constant ( $1 - 2$ ).
- ▶ PPMI ist ein measure for **first-order**-similarity, but we can also compute second-order similarity from  $A$  [2].



1. Fundamentals
2. Preprocessing and Heuristics
3. Statistical Word Similarity
4. Syntactical Word Similarity

# Syntactical Word Similarity: Edit Distances



## Now: Syntactical Word Similarity

- ▶ ... measures how similar two strings are in terms of their characters.
- ▶ useful for typos or variations of the same word (*above*).
- ▶ Common approach: **string edit distances**.




- ▶ Given two words/strings  $s_1, s_2$  of length  $n_1, n_2$ , an **edit distance** corresponds to the **number of operations** required to transform  $s_1$  into  $s_2$ .
- ▶ Different distances allow **different operations**:
  - ▶ **Levenshtein distance**:  
inserting/deleting/replacing 1 character.
  - ▶ **Damerau-Levenshtein distance**:  
also, transpose two adjacent characters.
  - ▶ ...

## Levenshtein Distance: Examples

- ▶  $D(\text{'typos'}, \text{'typohs'}) = 1$
- ▶  $D(\text{'Weisbahdn'}, \text{'Wiesbaden'}) = 4$



# Computing Levenshtein: Dynamic Programming



- ▶ Create a  $(n_1 \times n_2)$  **cost matrix**  $D$ .
- ▶ Initialize  $D$ 's first row (+column) with  $0, 1, \dots, n_1$  ( $0, 1, \dots, n_2$ ).
- ▶ **traverse**  $D$  **row-wise**, and every row from **left to right**.  
Compute each cell  $(i, j)$ :

$$D_{i,j} \leftarrow \min \left( \underbrace{1 + D_{i-1,j}}_{\text{delete}}, \underbrace{1 + D_{i,j-1}}_{\text{insert}}, \underbrace{1_{s_1(i) \neq s_2(j)} + D_{i-1,j-1}}_{\text{replace}} \right)$$

- ▶ Finally,  $D_{n_1, n_2}$  contains the Levenshtein distance.

# Levenshtein Distance: Speed




- ▶ **Complexity?**  $\rightarrow O(n_1 \cdot n_2)$  ☹️
- ▶ **Problem:** In practice, we want to compare a query  $q$  string/word with 100,000s of other words!
- ▶ **Example:** query=wiesbahden, find the most similar vocabulary word to correct the typo.

## Approach

Build **index structures** for fast term similarity matching:


1. Symspell (“reduction indexing”)
2. Min Hashing

- ▶ a method for fast string indexing: Symspell by Wolf Garbe<sup>4</sup>.
- ▶ Step 1: For each string  $s$  in the vocabulary: delete  $\Delta = 1, 2, \dots, K$  characters and store the resulting versions  $s'$  in a hash map.



<sup>4</sup><https://github.com/wolfgarbe/SymSpell> (C#),  
<https://github.com/mammothb/sympellpy> (Python)

- ▶ Step 2: Given a query  $q$ , perform the same delete operations, and lookup all reduced queries  $q'$  in the hash map.
- ▶ If  $q' = s'$ , we know that  $D(q, s) \leq \Delta(s, s') + \Delta(q, q')$ .
- ▶ In the example:  $D(\text{think}, \text{tank}) \leq \Delta(\text{think}, \text{tnk}) + \Delta(\text{tnk}, \text{think}) = 3$  (in fact,  $D(\text{think}, \text{tank}) = 2$ ).





## Remarks

- ▶ Huge efficiency gain, particularly because replaces and inserts are avoided (*think languages with huge alphabets like Chinese*).
- ▶ For each candidate  $s'$  we find, we can compute the exact distance  $D(q, s)$ .
- ▶ The reduced versions  $q'$  are iterated by increasing the number of deletes. Thus, we can stop once we cannot find a better version than the best match so far.



- [1] **Marti A. Hearst.**  
Automatic acquisition of hyponyms from large text corpora.  
In Proceedings of the 14th Conference on Computational Linguistics - Volume 2, COLING '92, pages 539–545,  
Stroudsburg, PA, USA, 1992. Association for Computational Linguistics.
  
- [2] **A. Islam and D. Inkpen.**  
Second order co-occurrence PMI for determining the semantic similarity of words.  
In Proc. LREC 2006, pages 1033–1038, 2006.
  
- [3] **Daniel Jurafsky and James H. Martin.**  
Speech and Language Processing: An Introduction to Natural Language Processing (3rd Edition Draft Chapters).  
2017.